

## ▼ 1) Load dataset

```
# Read the dataset into a pandas DataFrame
import pandas as pd
import os

# Update this path if needed (using raw strings to handle backslashes correctly)
if os.path.exists(r'/content/student-por.csv'):
    df = pd.read_csv(r'/content/student-por.csv', sep=';')
elif os.path.exists(r'/content/student-mat.csv'):
    df = pd.read_csv(r'/content/student-mat.csv', sep=';')
else:
    # Try to find any csv in working dir
    csvs = [f for f in os.listdir('.') if f.lower().endswith('.csv')]
    if csvs:
        df = pd.read_csv(csvs[0], sep=';')
    else:
        df = pd.DataFrame()
        print('No CSV found in working directory. Please upload or set the correct path.')

print('Data shape:', df.shape)
df.head()
```

Data shape: (649, 33)

	school	sex	age	address	famsize	Pstatus	Medu	Fedu	Mjob	Fjob	...	famrel	freetime	goout
0	GP	F	18	U	GT3	A	4	4	at_home	teacher	...	4	3	4
1	GP	F	17	U	GT3	T	1	1	at_home	other	...	5	3	3
2	GP	F	15	U	LE3	T	1	1	at_home	other	...	4	3	2
3	GP	F	15	U	GT3	T	4	2	health	services	...	3	2	2
4	GP	F	16	U	GT3	T	3	3	other	other	...	4	3	2

5 rows × 33 columns

## ▼ 2) Exploratory Data Analysis (EDA)

Inspect distributions, missing values, correlations, and class balance.

```
# Basic EDA
import matplotlib.pyplot as plt
import numpy as np
import seaborn as sns

# Set visualization style
sns.set_style('whitegrid')
plt.rcParams['figure.dpi'] = 100

print('\n--- Dataset Info ---')
display(df.info())

print('\n--- Descriptive Statistics (numeric) ---')
display(df.describe())

print('\n--- Missing Values ---')
```

```
missing = df.isnull().sum()
if missing.sum() > 0:
    display(missing[missing > 0])
else:
    print('No missing values found!')

# Plot distributions for numeric features (one chart per numeric column)
num_cols = df.select_dtypes(include=[np.number]).columns.tolist()
print(f'\nPlotting distributions for {len(num_cols)} numeric features...')
for col in num_cols[:15]: # Limit to first 15 to avoid too many plots
    plt.figure(figsize=(6,2.5))
    plt.hist(df[col].dropna(), bins=30, edgecolor='black', alpha=0.7)
    plt.title(f'Distribution: {col}')
    plt.xlabel(col)
    plt.ylabel('Count')
    plt.grid(True, alpha=0.3)
    plt.show()

# Plot categorical value counts for top categorical columns
cat_cols = df.select_dtypes(include=['object','category']).columns.tolist()
print(f'\nPlotting value counts for {min(len(cat_cols), 6)} categorical features...')
for col in cat_cols[:6]:
    plt.figure(figsize=(6,2.5))
    df[col].value_counts().plot(kind='bar', edgecolor='black', alpha=0.7)
    plt.title(f'Value Counts: {col}')
    plt.xlabel(col)
    plt.ylabel('Count')
    plt.xticks(rotation=45, ha='right')
    plt.grid(True, alpha=0.3, axis='y')
    plt.tight_layout()
    plt.show()
```



```
--- Dataset Info ---
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 649 entries, 0 to 648
Data columns (total 33 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   school      649 non-null    object  
 1   sex          649 non-null    object  
 2   age          649 non-null    int64  
 3   address     649 non-null    object  
 4   famsize     649 non-null    object  
 5   Pstatus      649 non-null    object  
 6   Medu         649 non-null    int64  
 7   Fedu         649 non-null    int64  
 8   Mjob          649 non-null    object  
 9   Fjob          649 non-null    object  
 10  reason        649 non-null    object  
 11  guardian     649 non-null    object  
 12  traveltimes  649 non-null    int64  
 13  studytime    649 non-null    int64  
 14  failures     649 non-null    int64  
 15  schoolsup    649 non-null    object  
 16  famsup       649 non-null    object  
 17  paid          649 non-null    object  
 18  activities   649 non-null    object  
 19  nursery       649 non-null    object  
 20  higher        649 non-null    object  
 21  internet     649 non-null    object  
 22  romantic     649 non-null    object  
 23  famrel        649 non-null    int64  
 24  freetime      649 non-null    int64  
 25  goout         649 non-null    int64  
 26  Dalc          649 non-null    int64  
 27  Walc          649 non-null    int64  
 28  health         649 non-null    int64  
 29  absences      649 non-null    int64  
 30  G1             649 non-null    int64  
 31  G2             649 non-null    int64  
 32  G3             649 non-null    int64  
dtypes: int64(16), object(17)
memory usage: 167.4+ KB
None
```

--- Descriptive Statistics (numeric) ---

	age	Medu	Fedu	traveltimes	studytime	failures	famrel	freetime
<b>count</b>	649.000000	649.000000	649.000000	649.000000	649.000000	649.000000	649.000000	649.000000
<b>mean</b>	16.744222	2.514638	2.306626	1.568567	1.930663	0.221880	3.930663	3.180277
<b>std</b>	1.218138	1.134552	1.099931	0.748660	0.829510	0.593235	0.955717	1.051093
<b>min</b>	15.000000	0.000000	0.000000	1.000000	1.000000	0.000000	1.000000	1.000000
<b>25%</b>	16.000000	2.000000	1.000000	1.000000	1.000000	0.000000	4.000000	3.000000
<b>50%</b>	17.000000	2.000000	2.000000	1.000000	2.000000	0.000000	4.000000	3.000000
<b>75%</b>	18.000000	4.000000	3.000000	2.000000	2.000000	0.000000	5.000000	4.000000
<b>max</b>	22.000000	4.000000	4.000000	4.000000	4.000000	3.000000	5.000000	5.000000

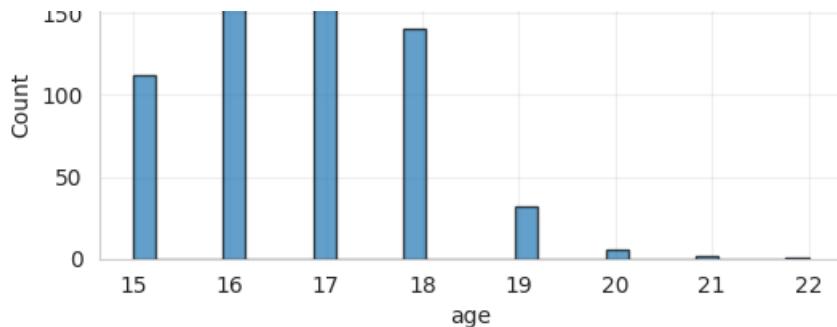
--- Missing Values ---

No missing values found!

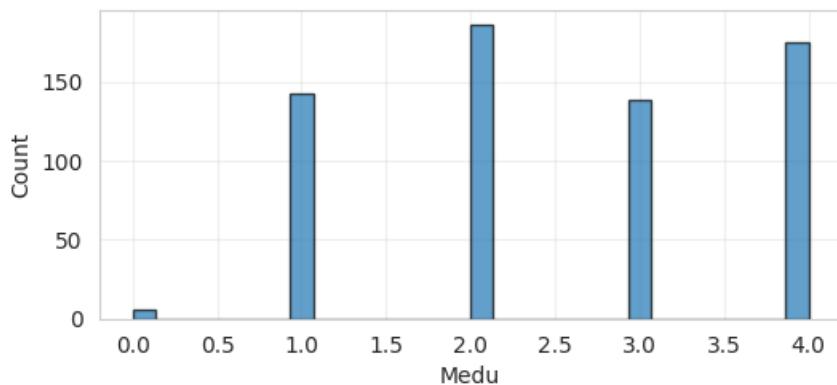
Plotting distributions for 16 numeric features...

Distribution: age

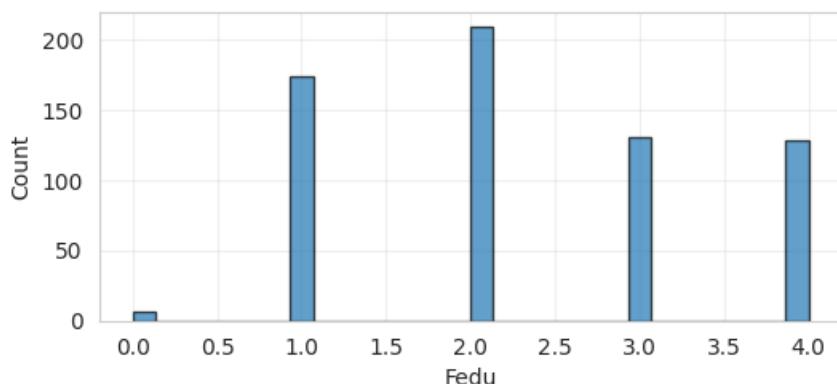




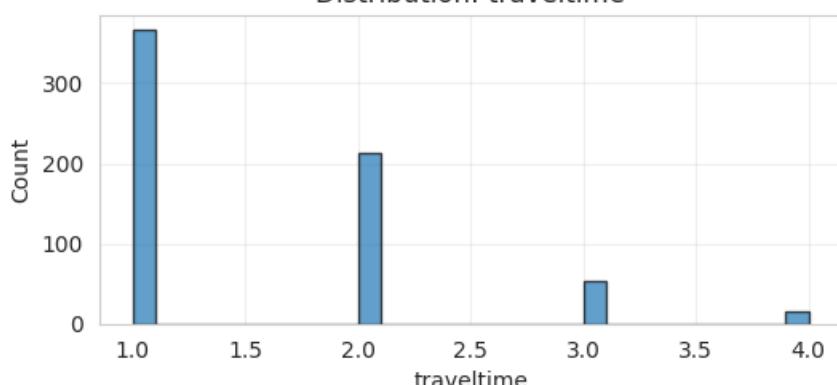
Distribution: Medu



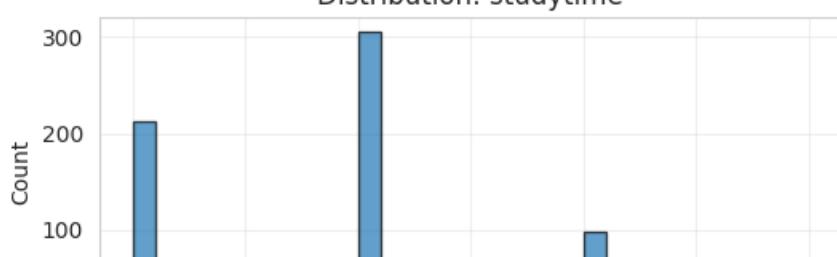
Distribution: Fedu



Distribution: traveltime

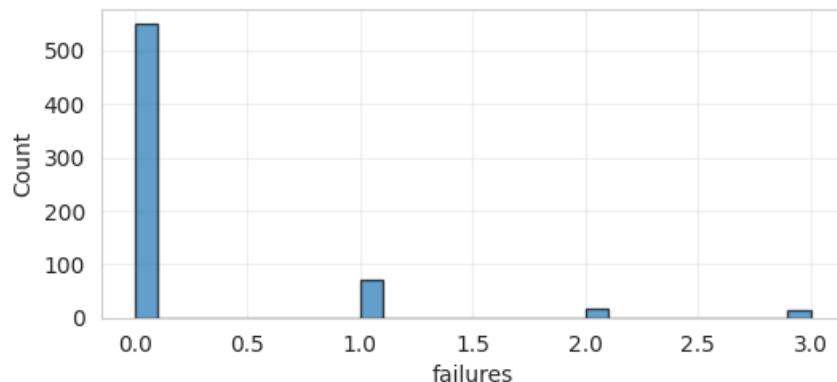


Distribution: studytime

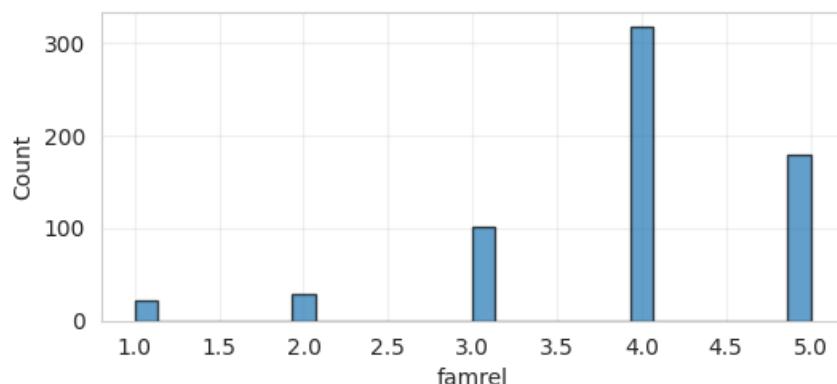




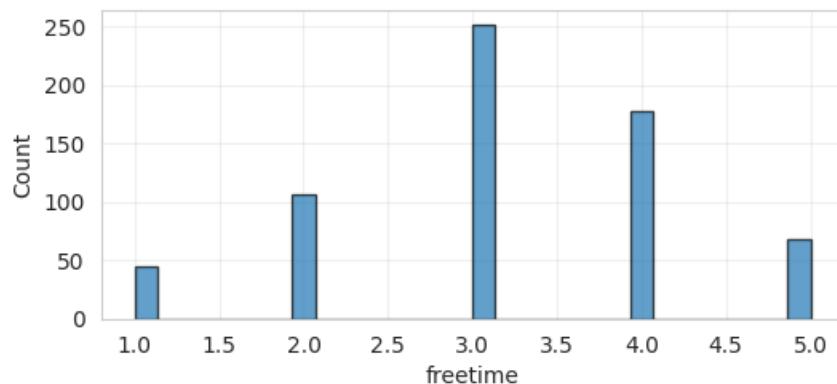
Distribution: failures



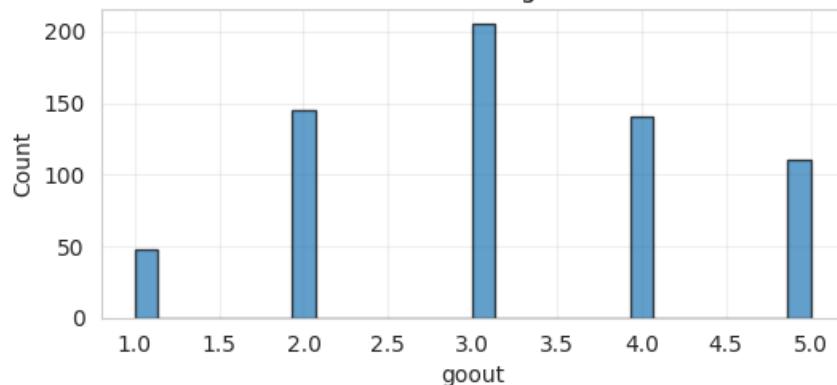
Distribution: famrel



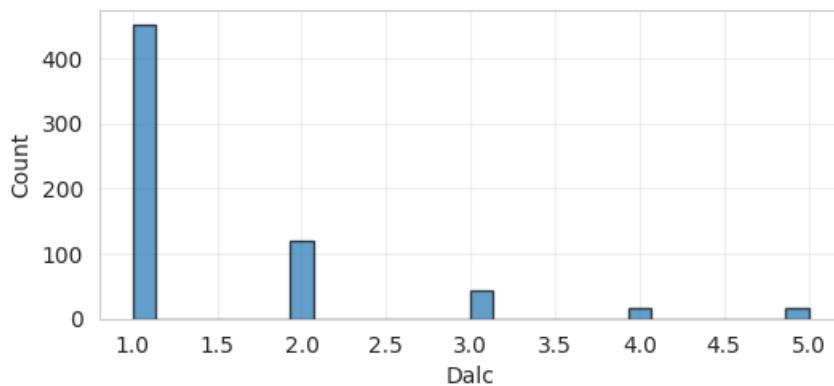
Distribution: freetime



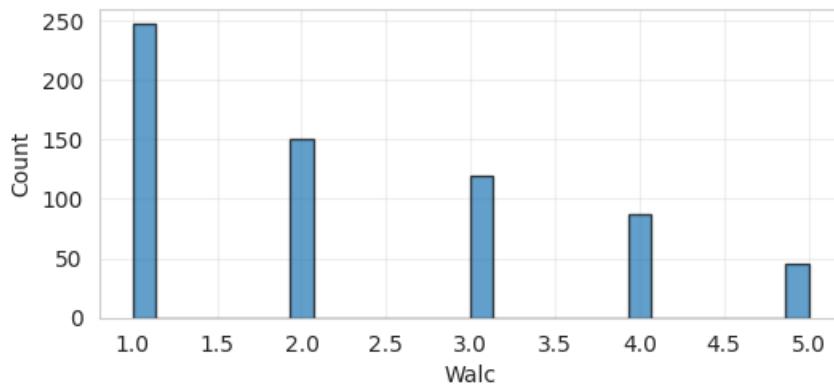
Distribution: goout



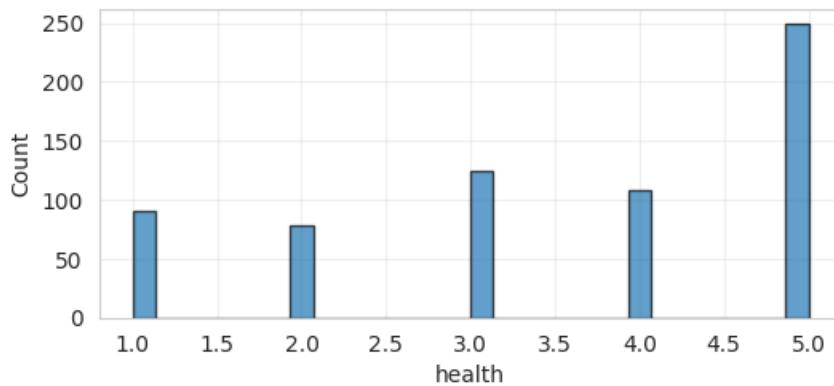
Distribution: Dalc



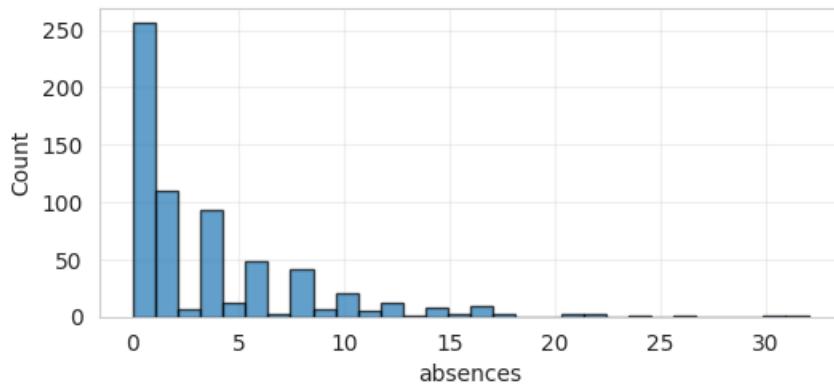
Distribution: Walc



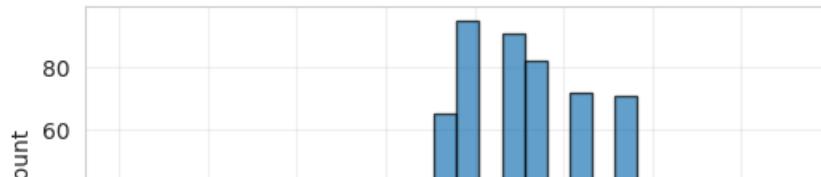
Distribution: health

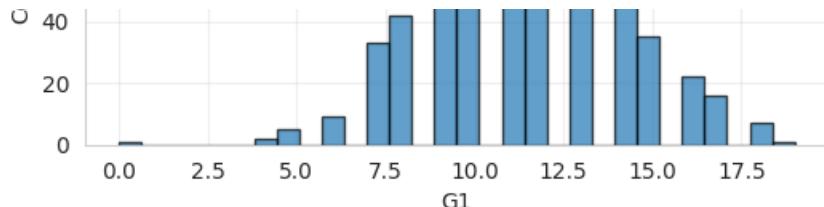


Distribution: absences

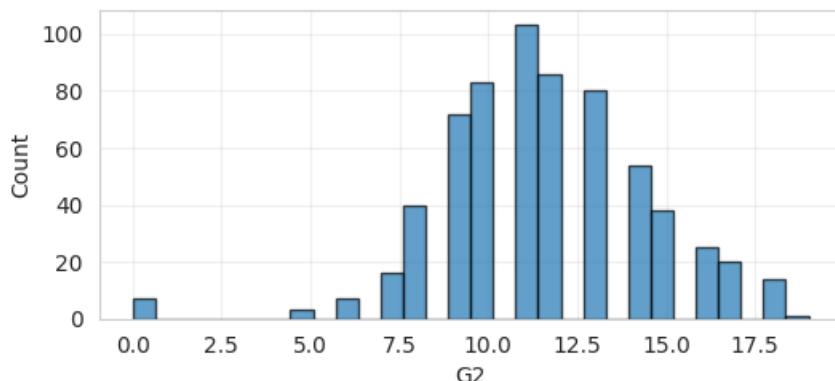


Distribution: G1



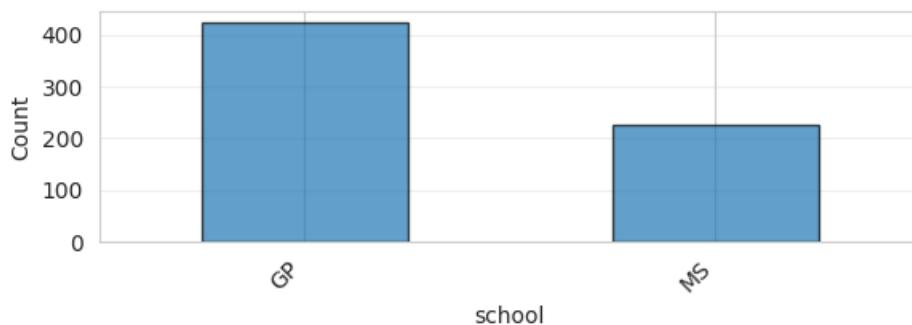


Distribution: G2

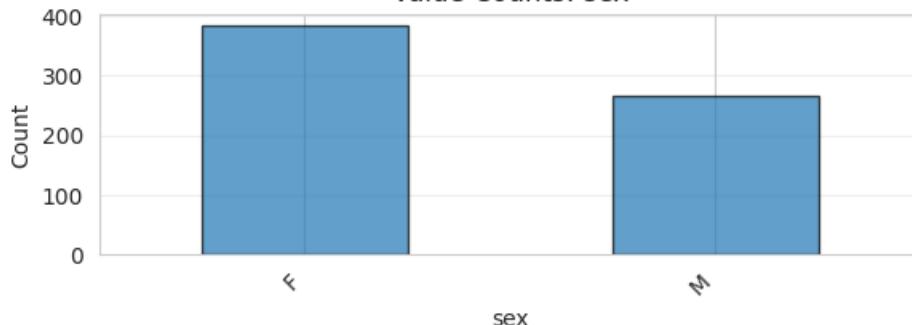


Plotting value counts for 6 categorical features...

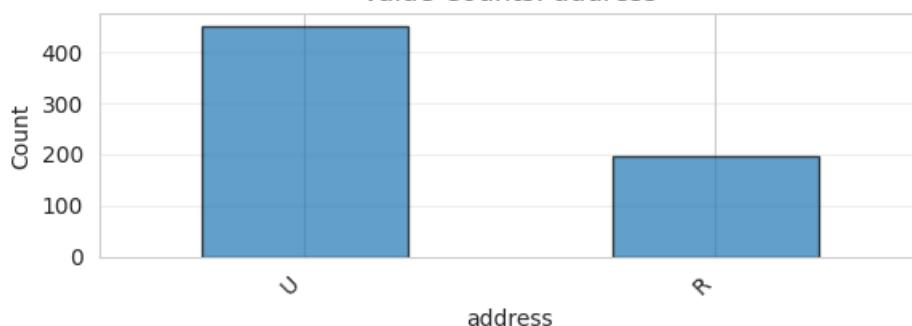
Value Counts: school



Value Counts: sex



Value Counts: address

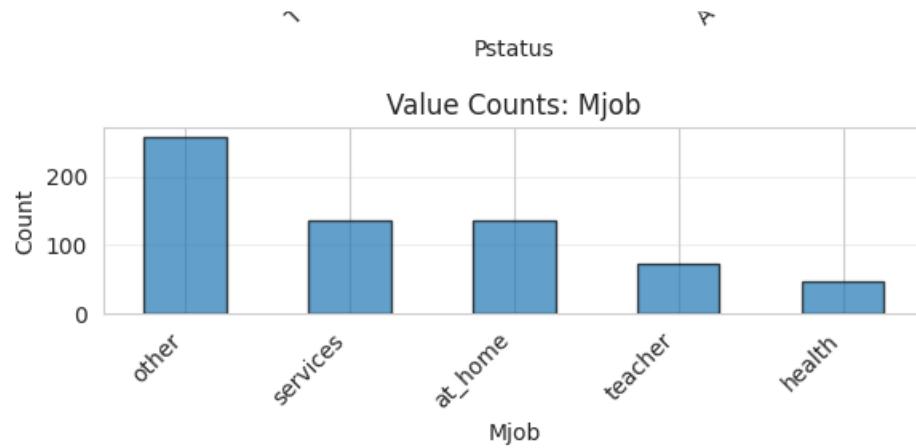


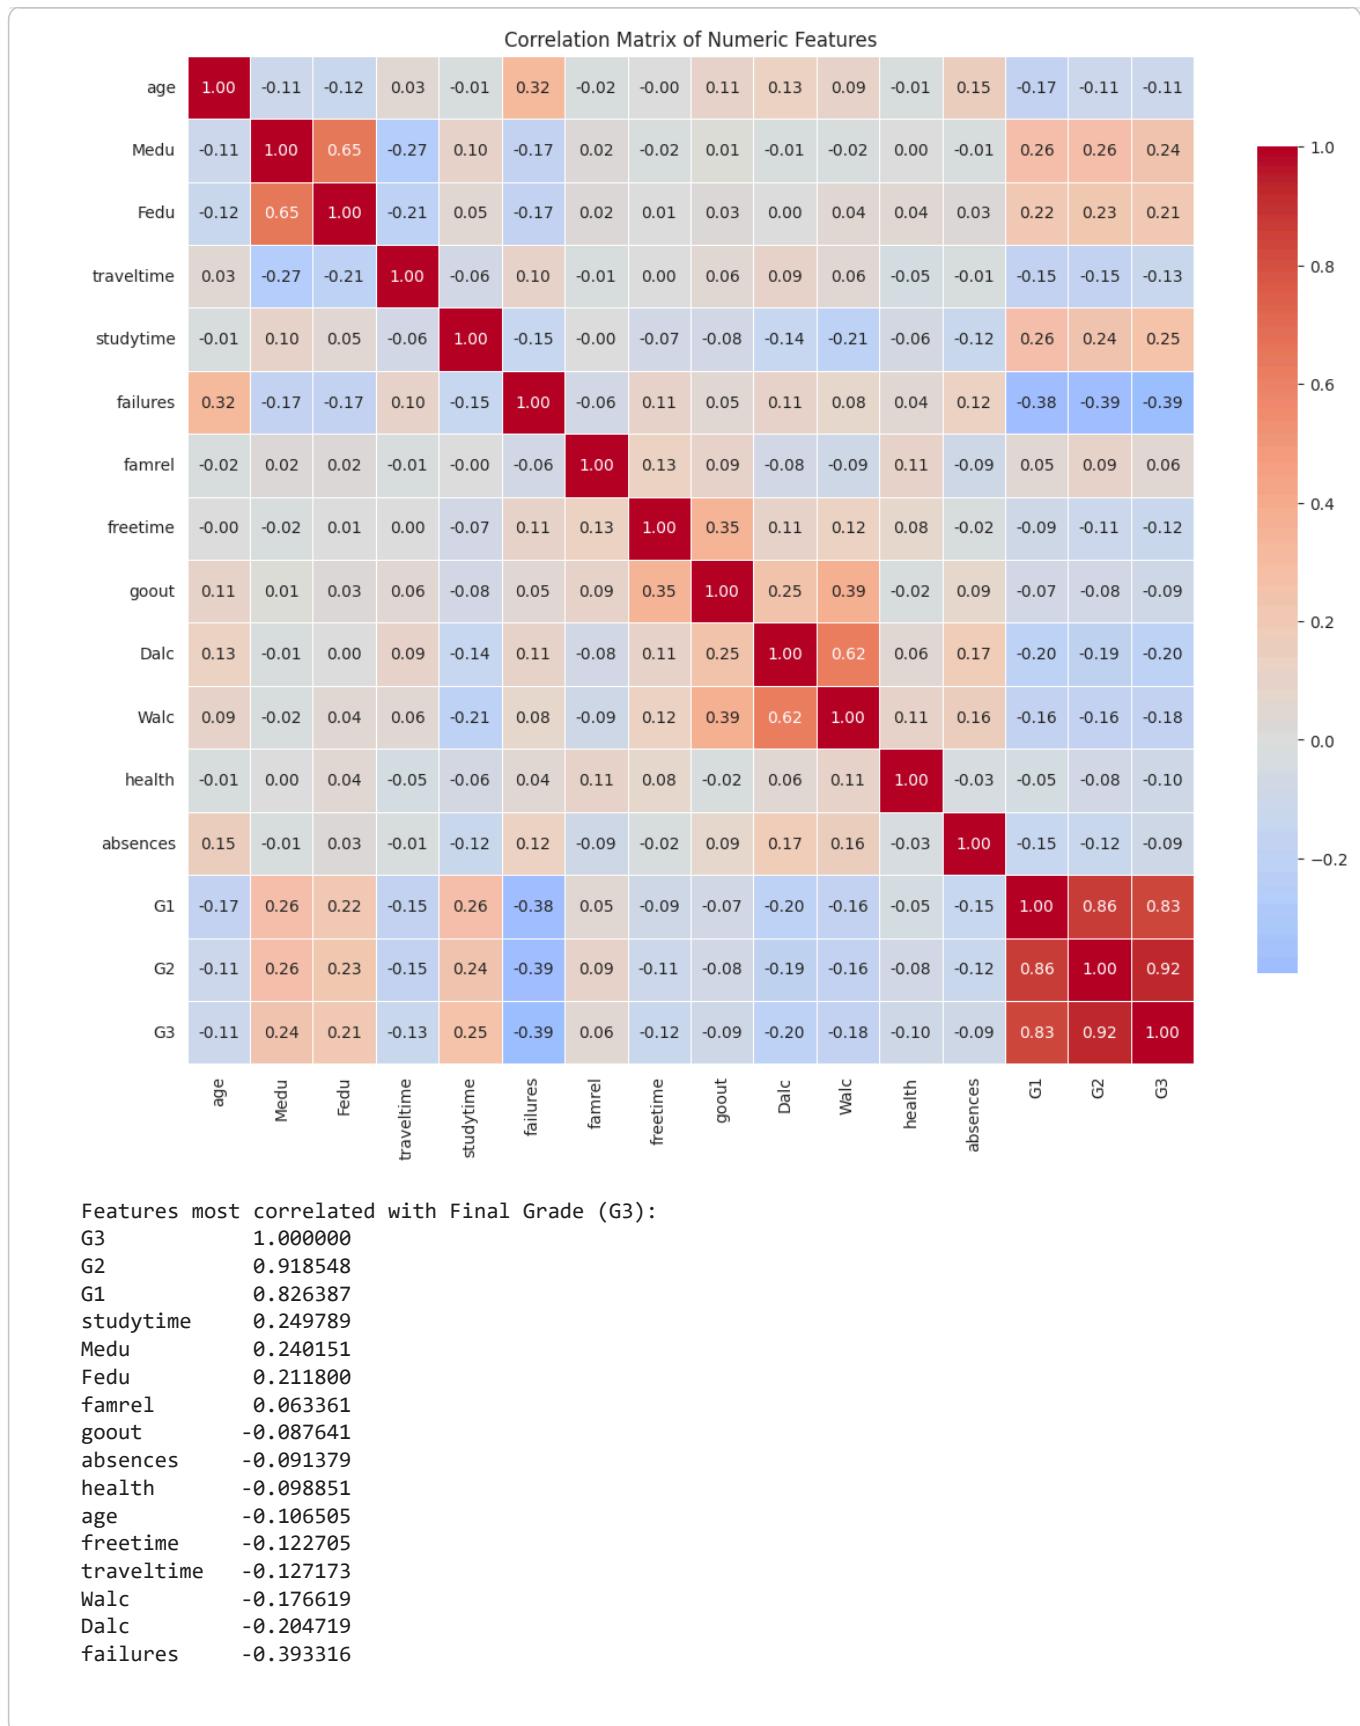
Value Counts: famsize



```
# Correlation analysis for numeric features
numeric_df = df.select_dtypes(include=[np.number])
if len(numeric_df.columns) > 1:
    plt.figure(figsize=(12, 10))
    correlation_matrix = numeric_df.corr()
    sns.heatmap(correlation_matrix, annot=True, fmt='.2f', cmap='coolwarm', center=0,
                square=True, linewidths=0.5, cbar_kws={"shrink": 0.8})
    plt.title('Correlation Matrix of Numeric Features')
    plt.tight_layout()
    plt.show()

# Show features most correlated with G3 (final grade)
if 'G3' in correlation_matrix.columns:
    print('\nFeatures most correlated with Final Grade (G3):')
    g3_corr = correlation_matrix['G3'].sort_values(ascending=False)
    print(g3_corr.to_string())
else:
    print('Not enough numeric columns for correlation analysis.')
```





Features most correlated with Final Grade (G3):

```

G3          1.000000
G2          0.918548
G1          0.826387
studytime   0.249789
Medu        0.240151
Fedu        0.211800
famrel      0.063361
goout       -0.087641
absences    -0.091379
health      -0.098851
age         -0.106505
freetime    -0.122705
traveltimes -0.127173
Walc        -0.176619
Dalc        -0.204719
failures    -0.393316

```

```
# Check class balance for target variable (G3 grades)
```

```
if 'G3' in df.columns:
```

```
    # Show distribution of final grades
    plt.figure(figsize=(10, 4))
```

```
    plt.subplot(1, 2, 1)
```

```

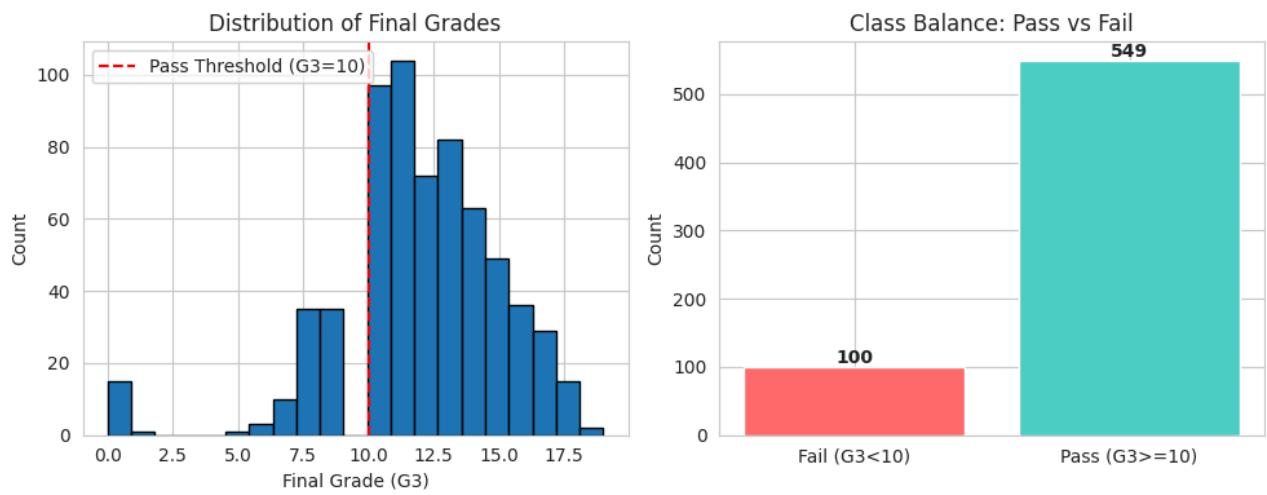
plt.hist(df['G3'], bins=21, edgecolor='black')
plt.xlabel('Final Grade (G3)')
plt.ylabel('Count')
plt.title('Distribution of Final Grades')
plt.axvline(x=10, color='r', linestyle='--', label='Pass Threshold (G3=10)')
plt.legend()

plt.subplot(1, 2, 2)
pass_count = (df['G3'] >= 10).sum()
fail_count = (df['G3'] < 10).sum()
plt.bar(['Fail (G3<10)', 'Pass (G3>=10)'], [fail_count, pass_count], color=['#ff6b6b', '#4ecdc4'])
plt.ylabel('Count')
plt.title('Class Balance: Pass vs Fail')
for i, v in enumerate([fail_count, pass_count]):
    plt.text(i, v + 5, str(v), ha='center', fontweight='bold')

plt.tight_layout()
plt.show()

print(f'Pass rate: {pass_count}/{len(df)} = {pass_count/len(df)*100:.1f}%')
print(f'Fail rate: {fail_count}/{len(df)} = {fail_count/len(df)*100:.1f}%')

```



Pass rate: 549/649 = 84.6%

Fail rate: 100/649 = 15.4%

### 3) Data Cleaning & Preprocessing

Handle missing values, encode categorical variables, and create target variable.

**Example target:** predict final grade pass/fail ( $G3 \geq 10 \rightarrow \text{pass}$ ). Adjust to your project's objective.

```

# Example preprocessing pipeline
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import OneHotEncoder, OrdinalEncoder
from sklearn.impute import SimpleImputer

# Make a copy
data = df.copy()

# Example: create binary target 'pass' from final grade column G3 if present
if 'G3' in data.columns:
    data['pass'] = (data['G3'] >= 10).astype(int)
    target_col = 'pass'

```

```

else:
    # If no numeric grade present, user should set target manually
    print('No G3 column found. Please define your target_col manually.')
    target_col = None

# Drop columns unlikely to be helpful (example: drop G1 and G2 to avoid leakage)
# G1 and G2 are period grades that directly predict G3, so excluding them makes the problem more realistic
drop_cols = ['G1', 'G2']
for c in drop_cols:
    if c in data.columns:
        data.drop(columns=c, inplace=True)

# Separate features and target
if target_col:
    X = data.drop(columns=[target_col])
    y = data[target_col]
else:
    X = data.copy()
    y = None

# Simple handling: numeric columns fillna with median, categorical with mode, one-hot encode categorical
num_cols = X.select_dtypes(include=[np.number]).columns.tolist()
cat_cols = X.select_dtypes(include=['object', 'category']).columns.tolist()

from sklearn.pipeline import Pipeline
from sklearn.compose import ColumnTransformer

num_pipe = Pipeline([('imputer', SimpleImputer(strategy='median'))])
cat_pipe = Pipeline([('imputer', SimpleImputer(strategy='most_frequent')),
                     ('onehot', OneHotEncoder(handle_unknown='ignore', sparse_output=False))])

preproc = ColumnTransformer([('num', num_pipe, num_cols),
                            ('cat', cat_pipe, cat_cols)], remainder='drop')

# Note: For supervised learning, preproc will be fit on training data only (see Cell 11)
# For unsupervised learning (clustering), we'll fit on the full dataset
print('Feature columns identified:')
print(f' - Numeric: {len(num_cols)} columns')
print(f' - Categorical: {len(cat_cols)} columns')

```

Feature columns identified:  
- Numeric: 14 columns  
- Categorical: 17 columns

```

# Extract column names after OneHotEncoder for interpretability
# Fit preproc on full X temporarily just to get feature names
try:
    preproc_temp = ColumnTransformer([('num', num_pipe, num_cols),
                                      ('cat', cat_pipe, cat_cols)], remainder='drop')
    preproc_temp.fit(X)

    ohe = None
    for name, trans, cols in preproc_temp.transformers_:
        if name == 'cat':
            ohe = trans.named_steps['onehot']
            cat_in_cols = cols
    feature_names = []
    # numeric names
    feature_names.extend(num_cols)
    # onehot names
    if ohe is not None:
        ohe_names = ohe.get_feature_names_out(cat_in_cols)
        feature_names.extend(list(ohe_names))
    print('Number of features after preprocessing:', len(feature_names))

```

```

except Exception as e:
    feature_names = None
    print('Could not extract feature names automatically.', e)

feature_names[:50] if feature_names else None

Number of features after preprocessing: 57
['age',
 'Medu',
 'Fedu',
 'traveltime',
 'studytime',
 'failures',
 'famrel',
 'freetime',
 'goout',
 'Dalc',
 'Walc',
 'health',
 'absences',
 'G3',
 'school_GP',
 'school_MS',
 'sex_F',
 'sex_M',
 'address_R',
 'address_U',
 'famsize_GT3',
 'famsize_LE3',
 'Pstatus_A',
 'Pstatus_T',
 'Mjob_at_home',
 'Mjob_health',
 'Mjob_other',
 'Mjob_services',
 'Mjob_teacher',
 'Fjob_at_home',
 'Fjob_health',
 'Fjob_other',
 'Fjob_services',
 'Fjob_teacher',
 'reason_course',
 'reason_home',
 'reason_other',
 'reason_reputation',
 'guardian_father',
 'guardian_mother',
 'guardian_other',
 'schoolsup_no',
 'schoolsup_yes',
 'famsup_no',
 'famsup_yes',
 'paid_no',
 'paid_yes',
 'activities_no',
 'activities_yes',
 'nursery_no']

```

## 4) Predictive Modeling (Supervised)

Three models: Logistic Regression, Random Forest, and Gradient Boosting.

```

# Modeling: train/test split and basic training pipeline
from sklearn.model_selection import train_test_split, cross_val_score, StratifiedKFold
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, roc_auc_score, confi

```

```
# Only proceed if we have a target
if y is None:
    print('No target variable defined. Define y to proceed with supervised modeling.')
else:
    X_train, X_test, y_train, y_test = train_test_split(X, y, stratify=y, test_size=0.2, random_state=42)
    # Fit preprocessor ONLY on training data to avoid data leakage
    X_train_pre = preproc.fit_transform(X_train)
    X_test_pre = preproc.transform(X_test)

    models = {
        'LogisticRegression': LogisticRegression(max_iter=1000, random_state=42),
        'RandomForest': RandomForestClassifier(n_estimators=200, random_state=42),
        'GradientBoosting': GradientBoostingClassifier(n_estimators=200, random_state=42)
    }

    results = {}
    for name, model in models.items():
        model.fit(X_train_pre, y_train)
        y_pred = model.predict(X_test_pre)
        y_proba = model.predict_proba(X_test_pre)[:,1] if hasattr(model, 'predict_proba') else None
        res = {
            'accuracy': accuracy_score(y_test, y_pred),
            'precision': precision_score(y_test, y_pred, zero_division=0),
            'recall': recall_score(y_test, y_pred, zero_division=0),
            'f1': f1_score(y_test, y_pred, zero_division=0),
        }
        if y_proba is not None:
            res['roc_auc'] = roc_auc_score(y_test, y_proba)
        results[name] = res
        print(f'--- {name} ---')
        print(classification_report(y_test, y_pred))
    print('\nSummary results:')
    import pandas as pd
    display(pd.DataFrame(results).T)
```

--- LogisticRegression ---				
	precision	recall	f1-score	support
0	1.00	0.85	0.92	20
1	0.97	1.00	0.99	110
accuracy			0.98	130
macro avg	0.99	0.93	0.95	130
weighted avg	0.98	0.98	0.98	130

```
# Visualize model performance comparison
if y is not None and len(results) > 0:
    import pandas as pd
    results_df = pd.DataFrame(results).T

    # Plot metrics comparison
    fig, axes = plt.subplots(2, 2, figsize=(12, 10))
    metrics = ['accuracy', 'precision', 'recall', 'f1']

    for idx, metric in enumerate(metrics):
        ax = axes[idx // 2, idx % 2]
        if metric in results_df.columns:
            results_df[metric].plot(kind='bar', ax=ax, color=['#3498db', '#2ecc71', '#e74c3c'])
            ax.set_title(f'{metric.capitalize()} by Model')
            ax.set_ylabel(metric.capitalize())
            ax.set_xlabel('Model')
            ax.set_ylim([0, 1])
            ax.grid(True, alpha=0.3, axis='y')
            ax.set_xticklabels(results_df.index, rotation=45, ha='right')

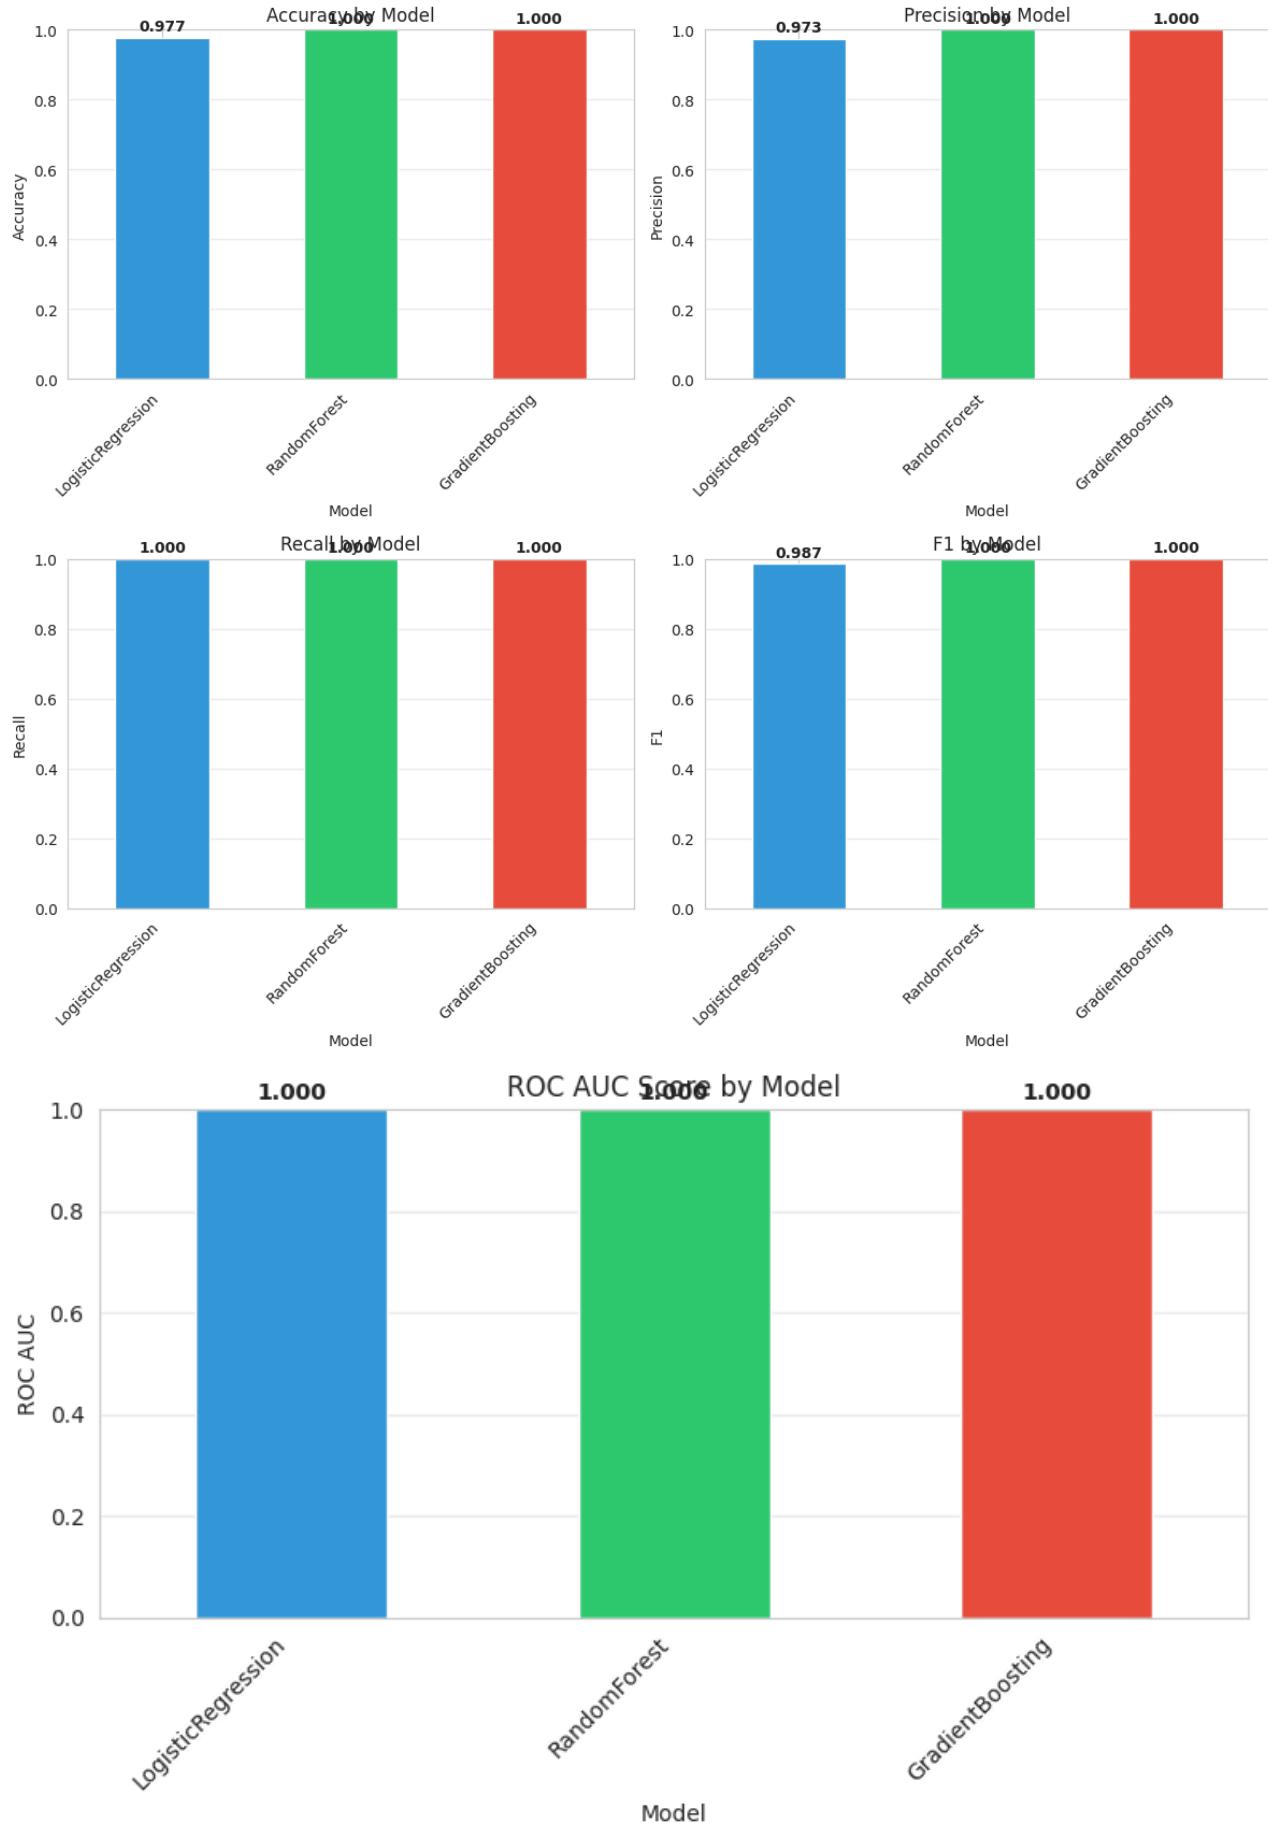
            # Add value labels on bars
            for i, v in enumerate(results_df[metric]):
                ax.text(i, v + 0.02, f'{v:.3f}', ha='center', fontweight='bold')

    plt.tight_layout()
    plt.show()

    # Display ROC AUC comparison if available
    if 'roc_auc' in results_df.columns:
        plt.figure(figsize=(8, 5))
        results_df['roc_auc'].plot(kind='bar', color=['#3498db', '#2ecc71', '#e74c3c'])
        plt.title('ROC AUC Score by Model')
        plt.ylabel('ROC AUC')
        plt.xlabel('Model')
        plt.ylim([0, 1])
        plt.grid(True, alpha=0.3, axis='y')
        plt.xticks(rotation=45, ha='right')

        # Add value labels
        for i, v in enumerate(results_df['roc_auc']):
            plt.text(i, v + 0.02, f'{v:.3f}', ha='center', fontweight='bold')

    plt.tight_layout()
    plt.show()
```



```
# Cross-validation (Stratified K-Fold) for each model
# Use Pipeline to avoid data leakage - preprocessor fits only on training folds
if y is not None:
    from sklearn.pipeline import Pipeline
    skf = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
    cv_results = {}

    # Create pipelines for each model to ensure proper cross-validation
    for name, model in models.items():
        # Create a fresh preprocessor for CV to avoid leakage
        from sklearn.compose import ColumnTransformer
        num_pipe_cv = Pipeline([('imputer', SimpleImputer(strategy='median'))])
        cat_pipe_cv = Pipeline([('imputer', SimpleImputer(strategy='most_frequent')),
                               ('onehot', OneHotEncoder(handle_unknown='ignore', sparse_output=False))])
        preproc_cv = ColumnTransformer([('num', num_pipe_cv, num_cols),
                                       ('cat', cat_pipe_cv, cat_cols)], remainder='drop')

        # Create pipeline with preprocessing and model
        pipeline = Pipeline([('preprocess', preproc_cv), ('classifier', model)])
        scores = cross_val_score(pipeline, X, y, cv=skf, scoring='f1')
        cv_results[name] = scores
print('Cross-val F1 scores:')
display(pd.DataFrame(cv_results))
```

Cross-val F1 scores:

	LogisticRegression	RandomForest	GradientBoosting	
0	1.000000	1.000000	1.0	!
1	0.995475	1.000000	1.0	
2	1.000000	1.000000	1.0	
3	1.000000	1.000000	1.0	
4	0.986425	0.995434	1.0	

## Feature importance (for tree-based models)

Show top features from RandomForest.

```
# Feature importances from RandomForest
import numpy as np
if 'RandomForest' in models and feature_names is not None:
    rf = models['RandomForest']
    importances = rf.feature_importances_
    idx = np.argsort(importances)[::-1]

    print('Top 20 Most Important Features:')
    print('-' * 50)
    for i, feat_idx in enumerate(idx[:20], 1):
        print(f'{i:2d}. {feature_names[feat_idx]:40s} {importances[feat_idx]:.4f}')

    # Visualize top features
    plt.figure(figsize=(10, 6))
    top_n = 15
    top_idx = idx[:top_n]
    plt.barh(range(top_n), importances[top_idx])
    plt.yticks(range(top_n), [feature_names[i] for i in top_idx])
    plt.xlabel('Feature Importance')
    plt.title('Top 15 Feature Importances (Random Forest)')
    plt.gca().invert_yaxis()
    plt.tight_layout()
```